

Product Engineering Intern (7 positions)

Overview & Role Specifications

Function:	Product Engineering: AI, agentic AI and cybersecurity
Degree:	B.E. or B.Tech in CSE, IT, AI and ML, ECE or a related branch; MCA; M.Sc in Computer Science, Data Science or Cyber Security
Converts to:	Product Engineer
Working Hours:	Core hours are 11:00 to 17:00 IST
Access:	Development, test and lab environments with synthetic data.

About the Role

You will build AI features, AI agents and security capabilities in our products and take them from design to release. The software runs in the cloud and on customer premises.

What You Will Do

- Build LLM features: retrieval over documents (RAG), structured extraction and evaluation.
- Build AI agents that plan, call tools through MCP and ask a person for approval.
- Red-team AI features and agents for prompt injection and data leakage. Build guardrails.
- Test web applications and APIs against the OWASP Top 10 security framework and fix what you find.
- Build security engineering features: asset discovery, vulnerability data pipelines and risk scoring.
- Write backend services in Python or TypeScript on PostgreSQL database, with screens in React.
- Package what you build in containers for the cloud and for on-premises installation.
- Write tests and evaluations for everything you ship. Give and receive code review.

Requirements & Selection Criteria

What You Must Bring:

- Python or TypeScript, with at least one project we can read on GitHub or run.
- How LLMs work in practice: tokens, context windows, embeddings and tool calling.
- Web and API fundamentals: HTTP, REST, JSON, authentication and access control.
- Security fundamentals: the OWASP Top 10 security framework and basic networking.
- SQL, Git and the Linux command line.

Good to Have:

- A RAG or agent project built with a framework such as LangGraph or a model provider's SDK.
- Security practice through CTFs, PortSwigger Web Security Academy, TryHackMe or Hack The Box.
- Docker, a cloud platform or open-weight models run with Ollama or vLLM.

Role Exercise:

Choose one task. Build a question-answering service over PDF documents we provide and score it on ten test questions. Or build an agent that checks DNS records and TLS certificate expiry for a list of domains. Or find and write up the vulnerabilities in a vulnerable web application we provide.